

Distributed Execution of Unikernel Applications on Container Orchestrator Platform Kubernetes for IoT Scenarios

Author: **Atakan Yenel**

Supervisor: Prof. Dr. Michael Gerndt

Advisor: Vladimir Podolskiy

Introduction

- Run very small virtual machines on hypervisors and IoT devices
- Orchestrate them from the cloud
- Use them side by side with the cloud ecosystem

Motivation

- Security
- Resource utilization
- Scaling time
- IoT scenarios

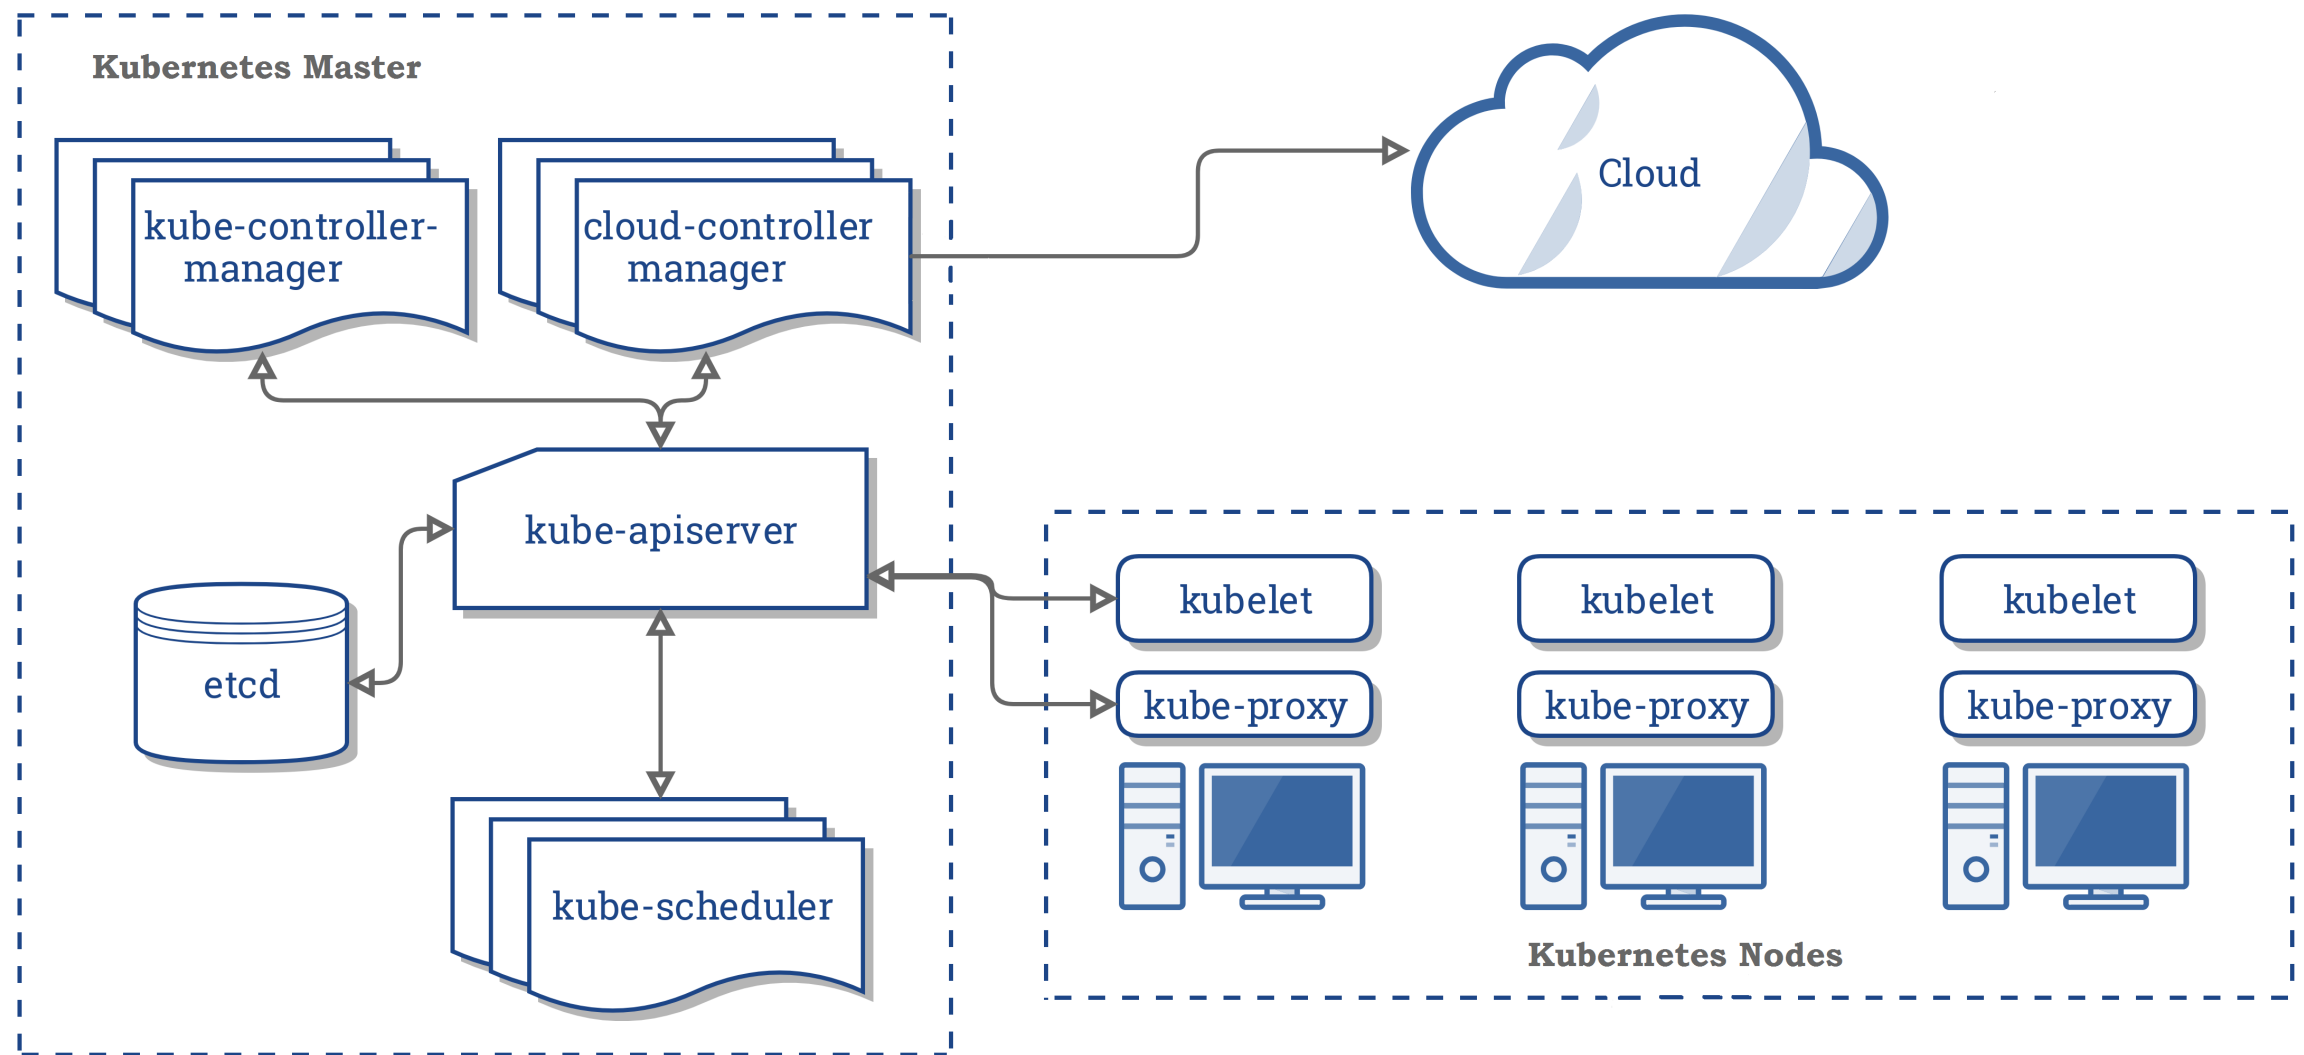
Background

- Kubernetes
- Unikernel

Kubernetes

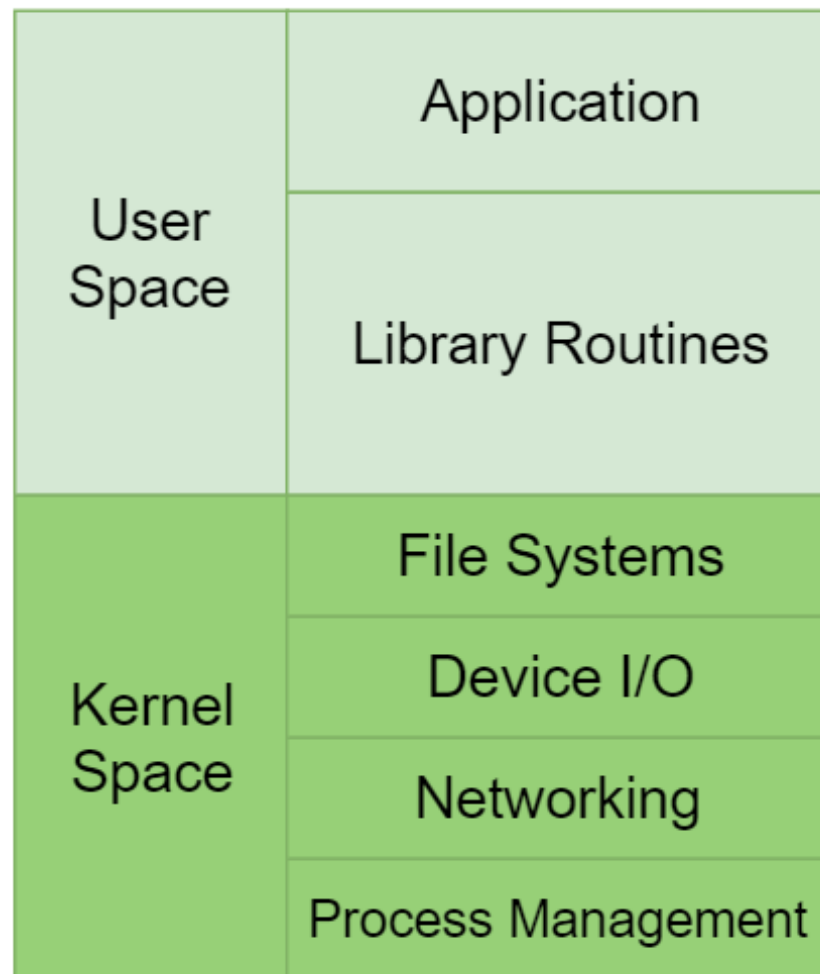
- Manages containerized applications across multiple hosts
- Developed & open sourced by Google
- Deployment & maintenance & **scaling** of applications
- Smallest deployable unit is pod, a group of containers
- Allows labeling hosts (nodes)

Kubernetes Architecture




Unikernel

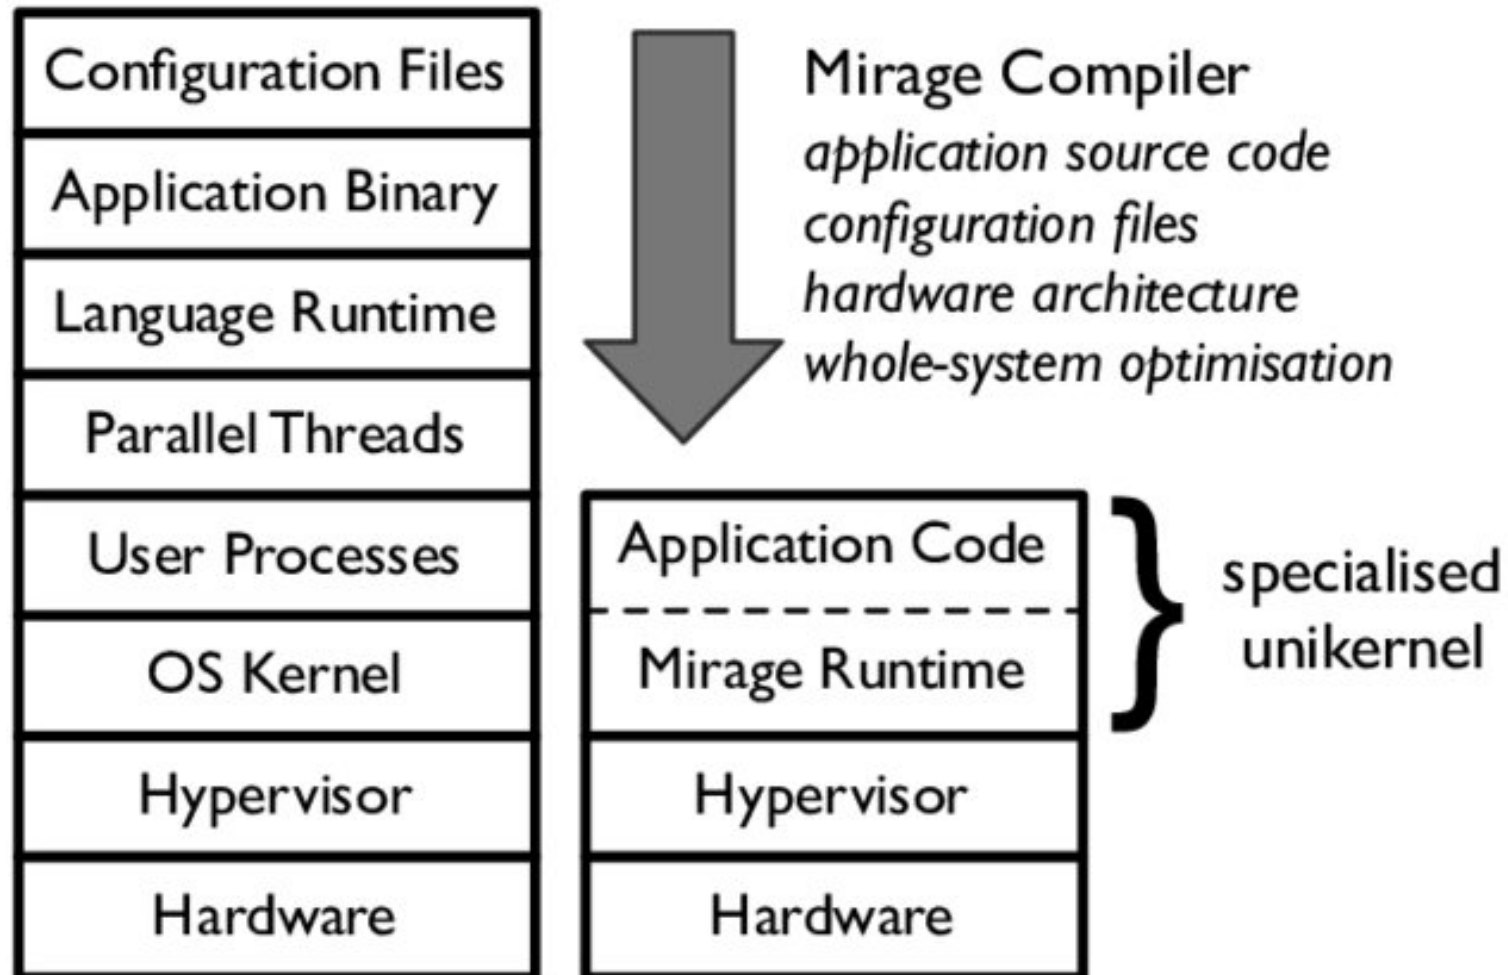
- Specialized, single address space machine images
- Minimal operating system for running the application code
- Runs directly on hypervisors & hardware
- First such system is Exokernel



Normal Application Stack



Unikernel Application Stack



Mirage OS

- Developed with OCaml
- From University of Cambridge & open sourced
- Runs on Xen & KVM
- ``mirage configure -t xen && make depend && mirage build``
- Provides libraries for:
 - Networking
 - Storage
 - Concurrency Support
- When compiled, libraries become OS drivers

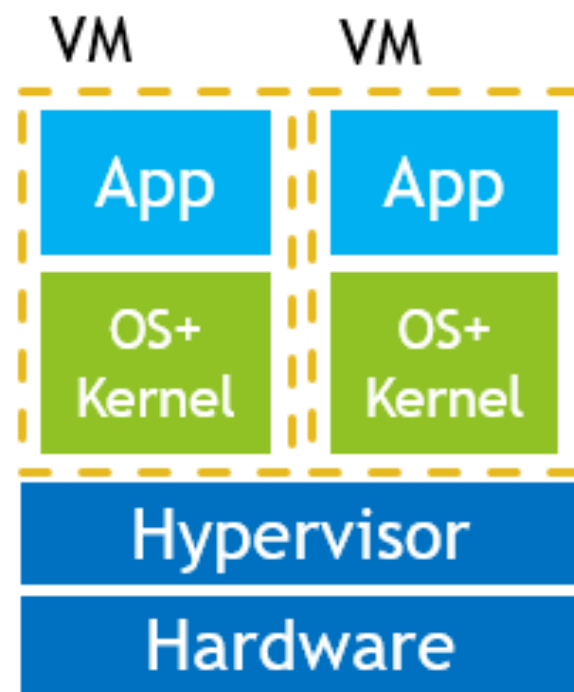
Mirage OS

```
open Lwt.Infix


module Main (KV: Mirage_kv.R0) (Time: Mirage_time.S) = struct

  let start kv _time=


    let read_from_file kv filename =
      KV.get kv (Mirage_kv.Key.v filename) >|= function
        | Error e ->
          Logs.warn (fun f -> f "Cannot find the file %a"
            KV.pp_error e)
        | Ok sensor_value ->
          Logs.info (fun f -> f "Reading from: %s -> %s" filename sensor_value);
    in
      let filename=Key_gen.filename() in
      let rec loop() =
        read_from_file kv filename >>= fun()->
        Time.sleep_ns (Duration.of_sec 2)>>= fun () ->
        loop()
      in
        loop()
  end
end
```



Virtual Machines




Linux Containers

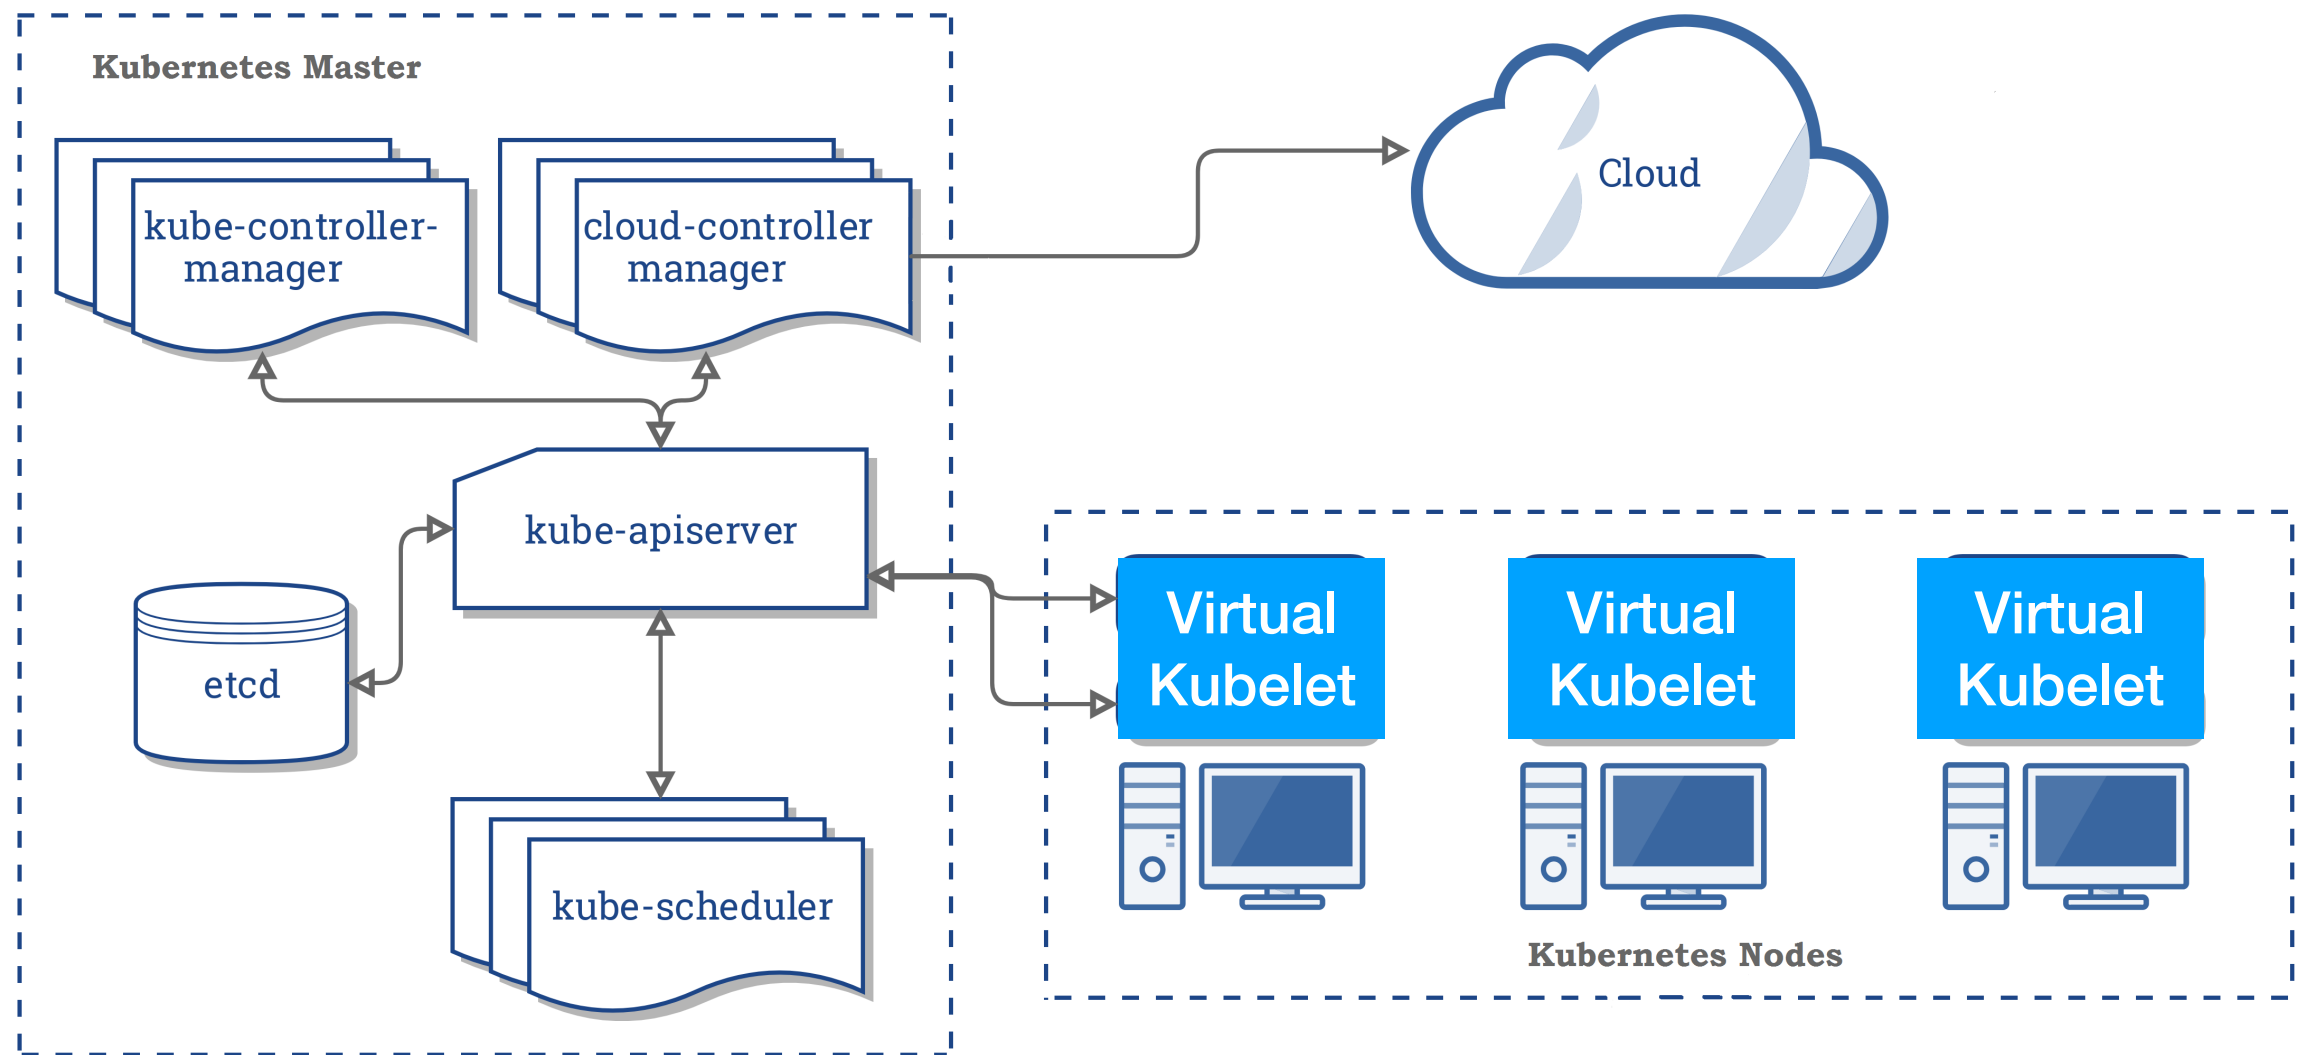


Unikernels

Implementation

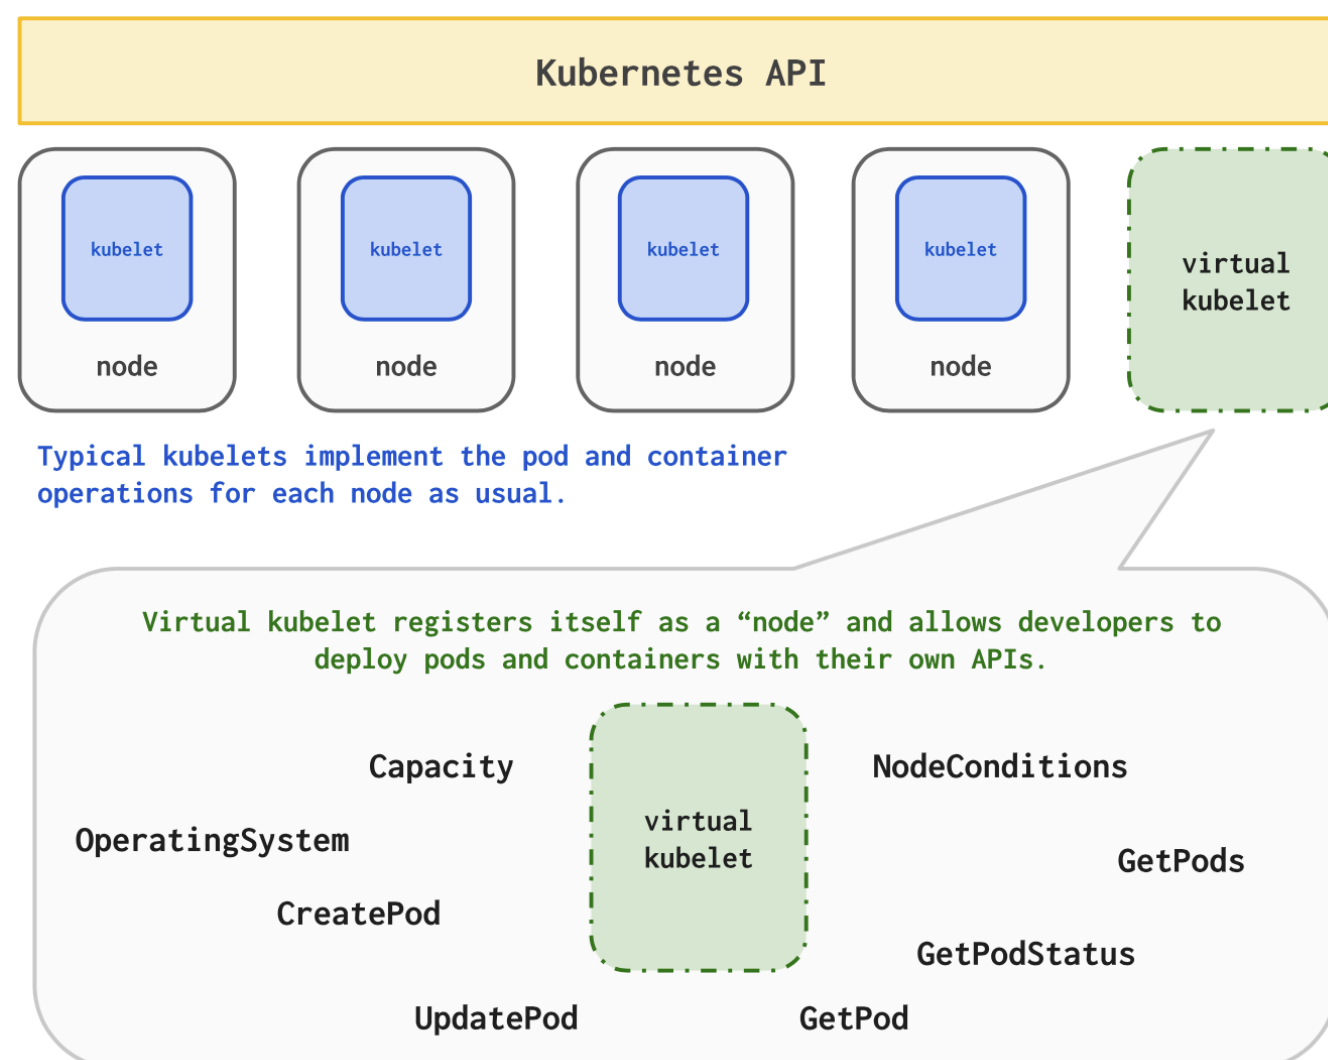


Implementation



Virtual-Kubelet

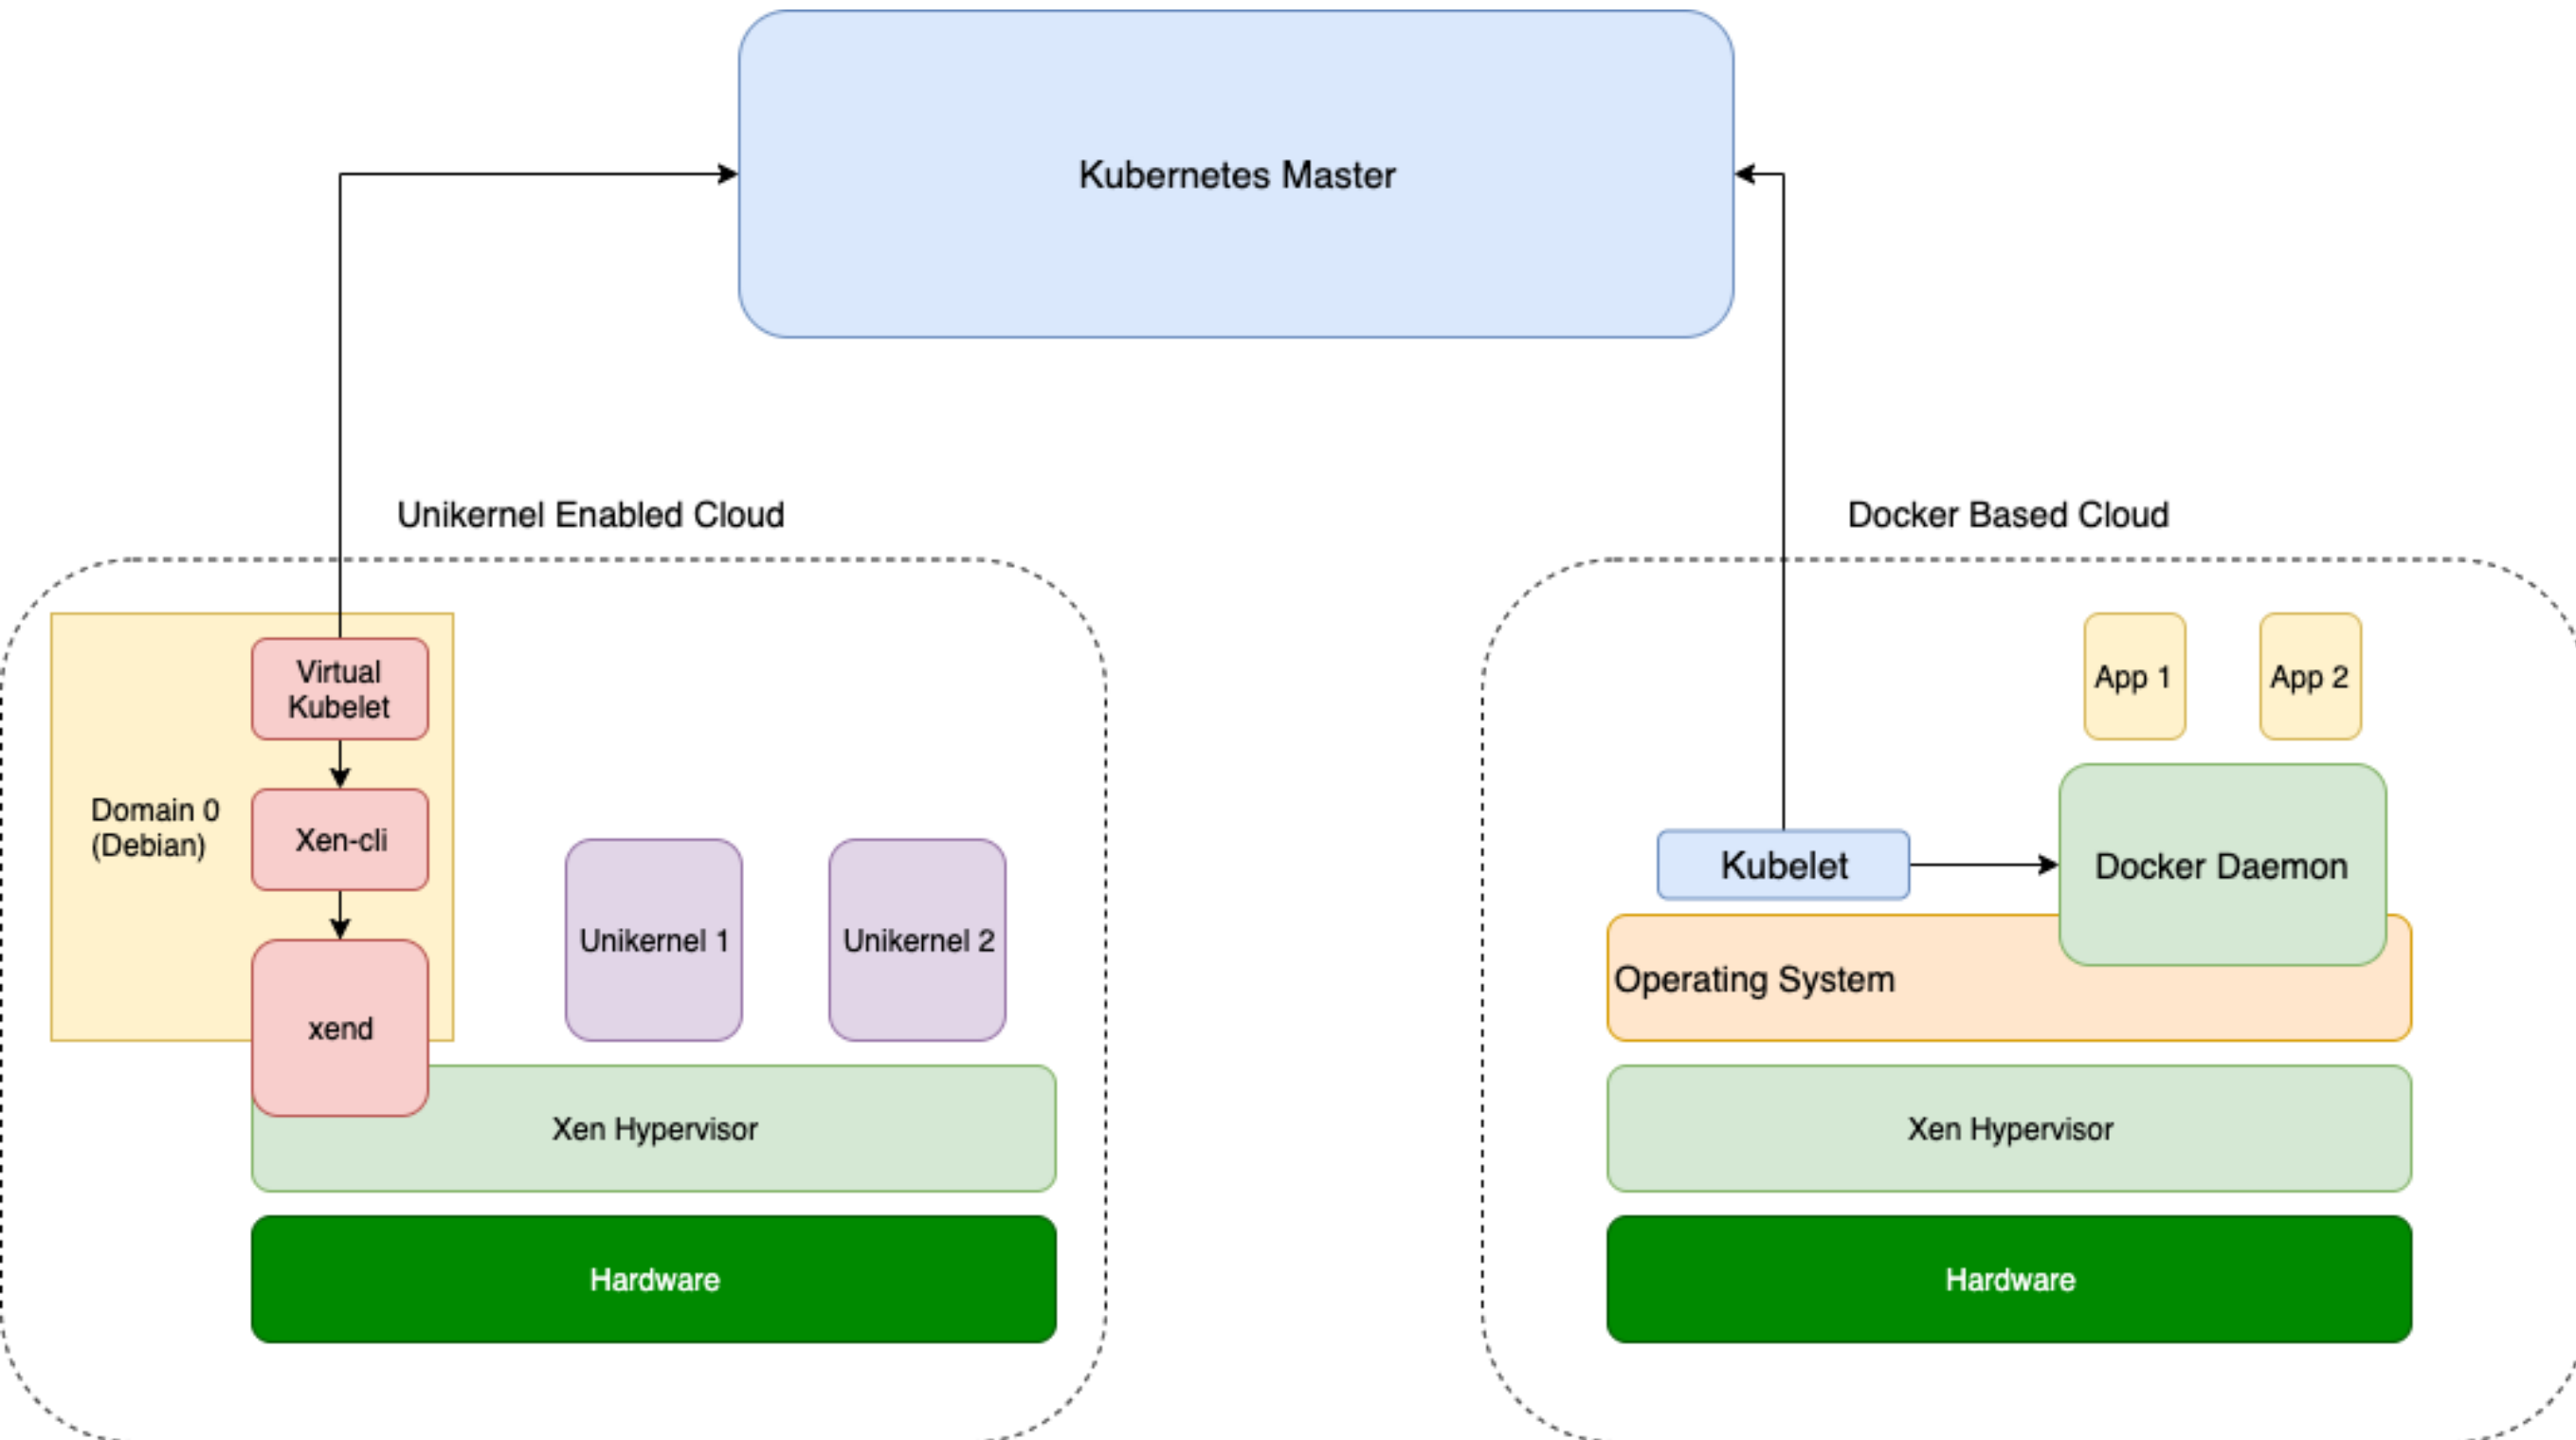
- A Kubelet Facade
- Developed by Microsoft
- Has an extensible API
- 36 MB for linux



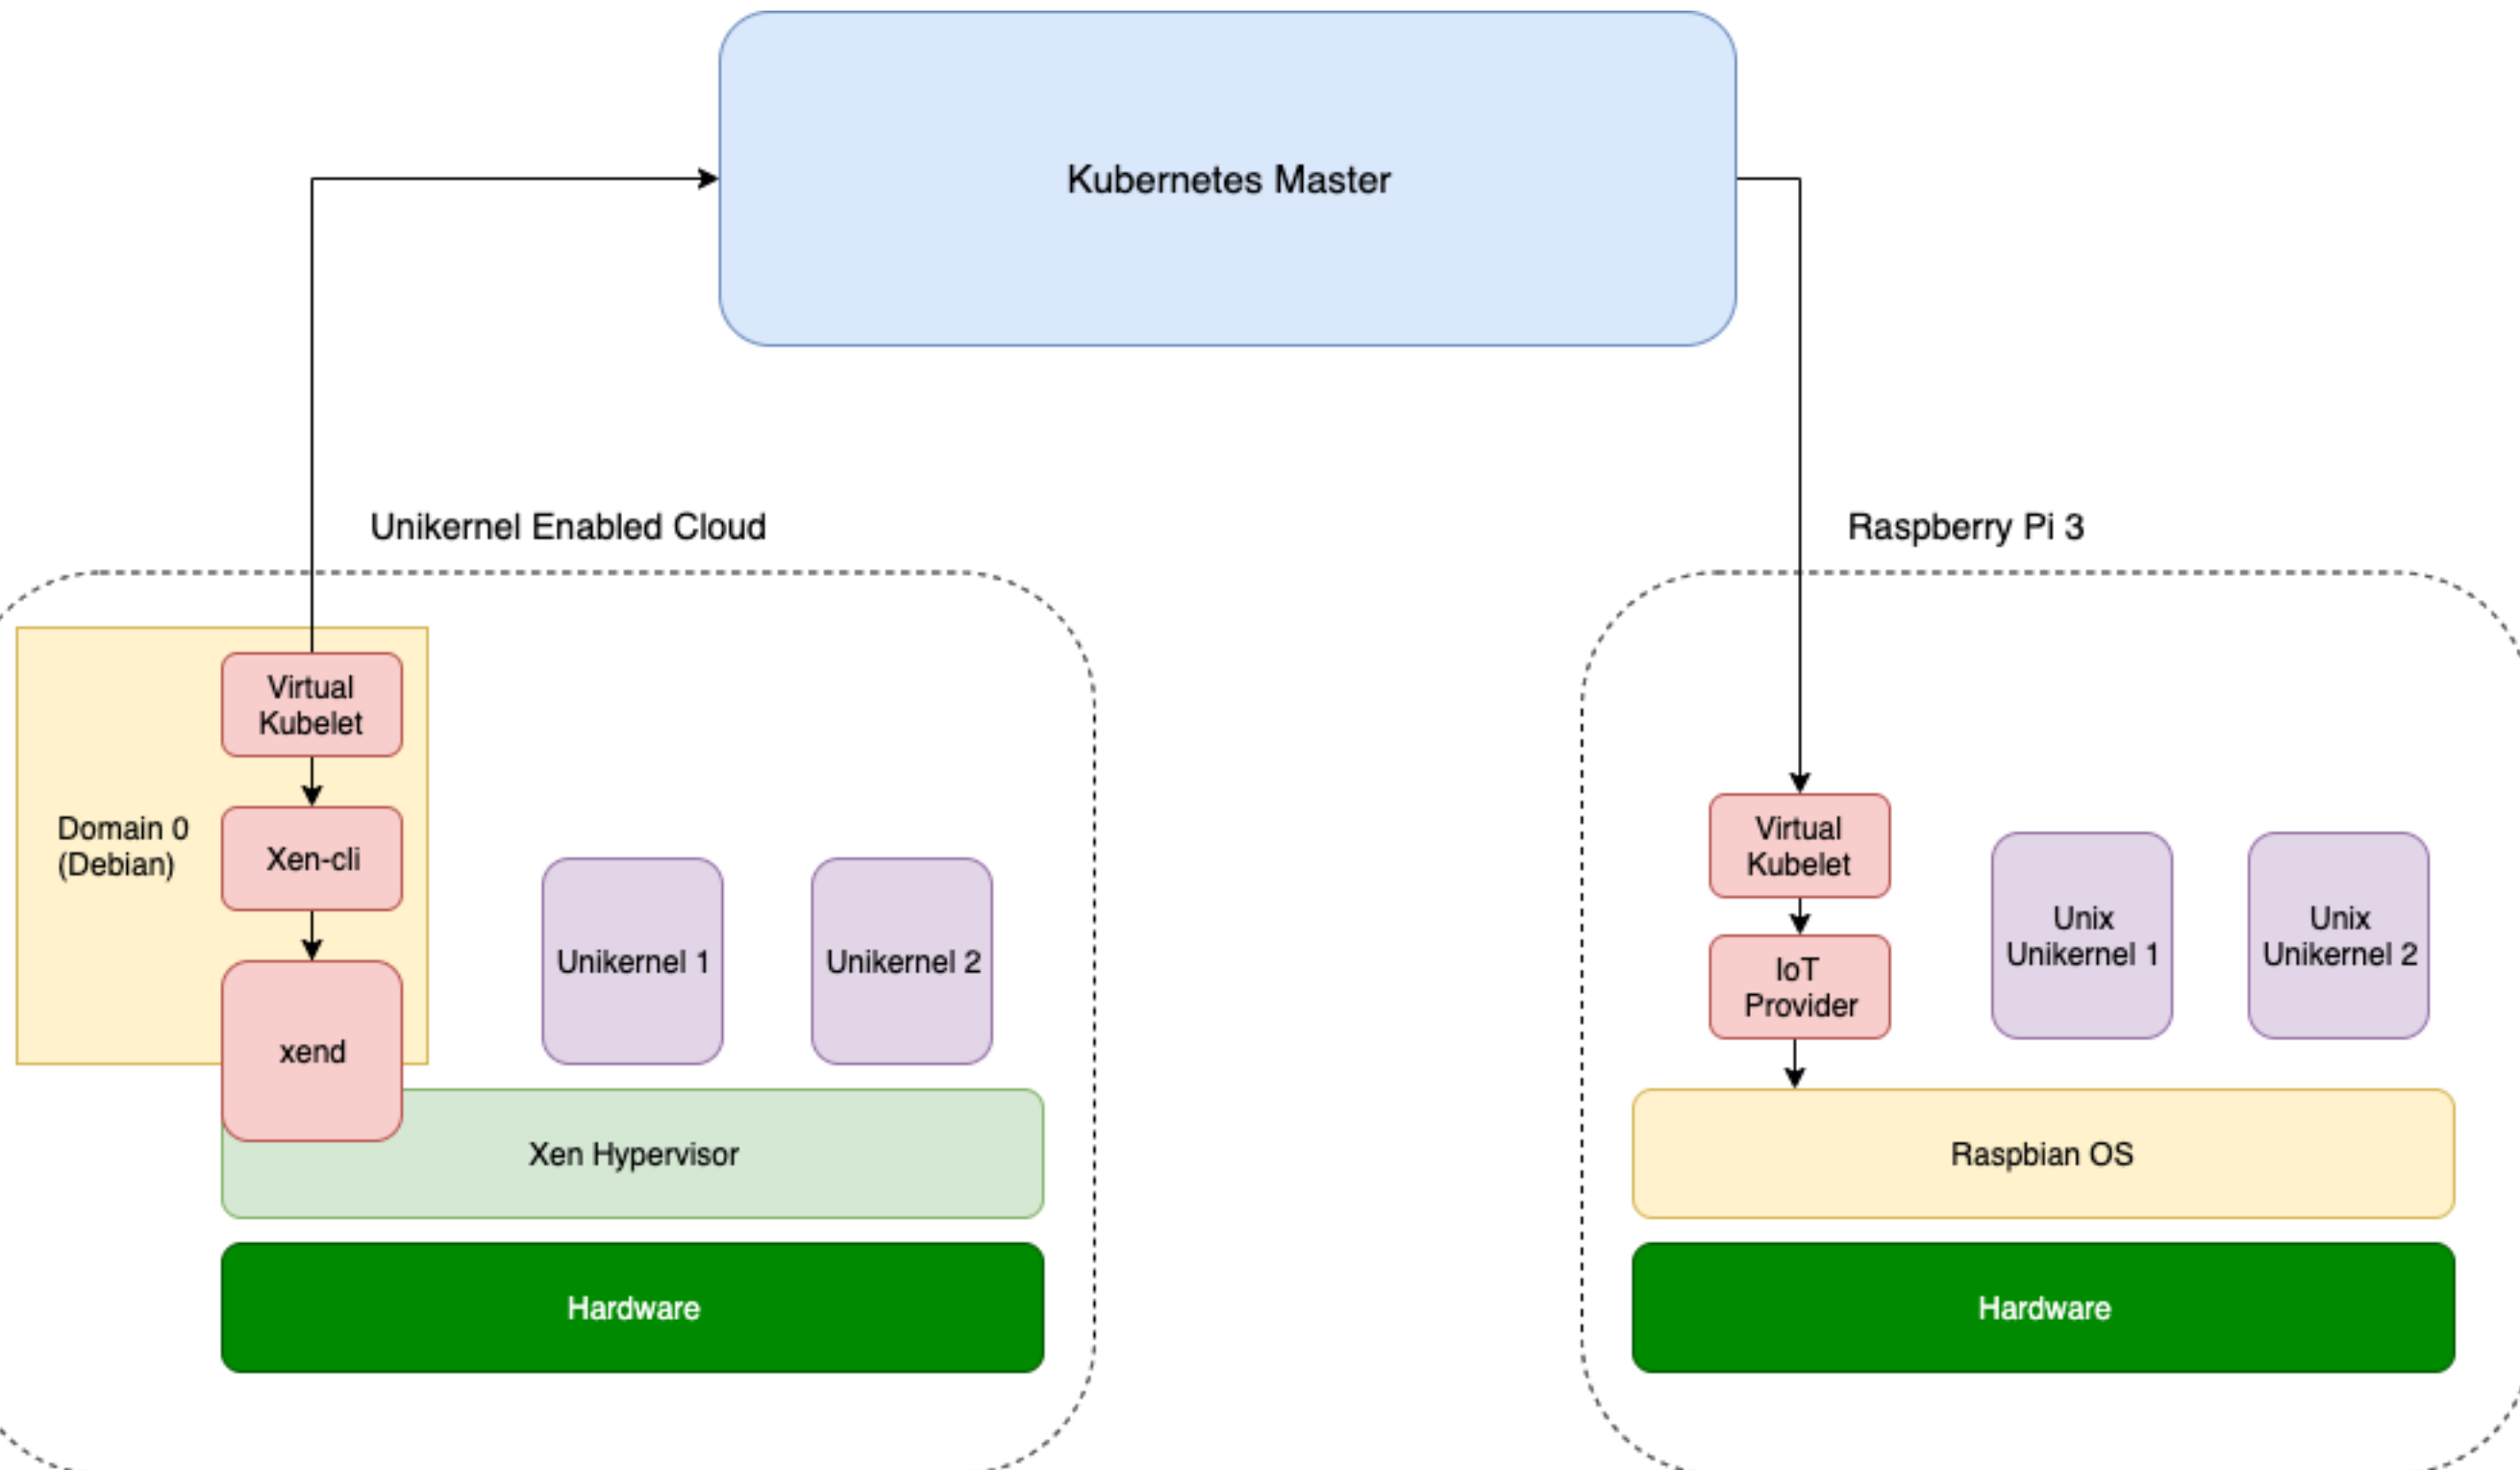
Kubelet API

```
type PodLifecycleHandler interface {  
  
    CreatePod(ctx context.Context, pod *corev1.Pod) error  
  
    UpdatePod(ctx context.Context, pod *corev1.Pod) error  
  
    DeletePod(ctx context.Context, pod *corev1.Pod) error  
  
    GetPod(ctx context.Context, namespace, name string) (*corev1.Pod, error)  
  
    GetPodStatus(ctx context.Context, namespace, name string) (*corev1.PodStatus, error)  
  
    GetPods(context.Context) ([]*corev1.Pod, error)  
  
}
```


Architecture



Managing IoT

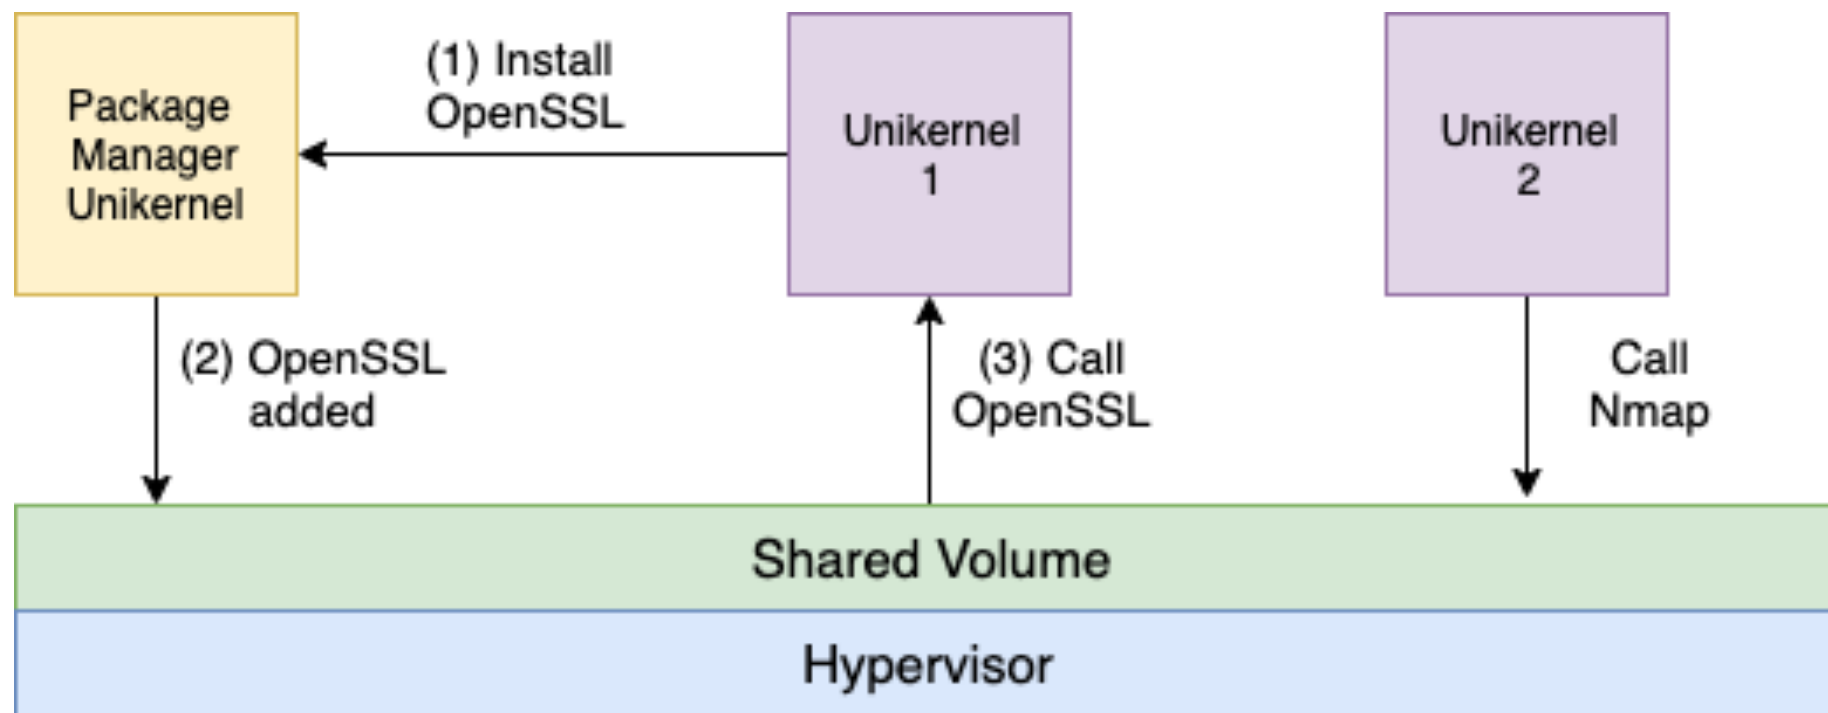


Label-based deployment

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: green-daemon
  namespace: default

spec:
  (***)
  spec:
    containers:
      - name: uni-run
        image: humidity-daemon
        args:
          - GREEN
    nodeSelector:
      type: virtual-kubelet
      location: germany
      sensor: humidity
      os: raspbian
    tolerations:
      - key: virtual-kubelet.io/provider
        operator: Equal
        value: unikernel
        effect: NoSchedule
```

Calling external packages




Evaluation

- Image Size
- Boot up
- Scalability
- Usability

Target	Size
Docker	69.4 MB
Unix	4.7 MB
Xen	3.3 MB
Virtio	2.4 MB
Hvt	2.3 MB

Scalability






Figure 5.5.: Scaling from 0 to 30 pods




Figure 5.6.: Scaling from 0 to 50 pods




Figure 5.7.: Scaling from 50 to 0 pods




Figure 5.8.: Scaling from 30 to 0 pods

Motivation revisited

- **Security**
 - Single address space
 - Only process is application
 - Isolated by hypervisor
- **Resource utilization**
 - Runs on hypervisor
 - No background processes
- **Scaling Time**
 - Faster than Docker
 - Smaller image size
- **IoT Scenarios**
 - Single interface for cloud & IoT
 - DaemonSet support
 - Kubernetes

Drawbacks

- Very limited language support for Unikernels
- Not all IoT devices support virtualization
- Kubernetes uses HTTP, additional overhead for IoT
- Docker is good enough
 - Has dependencies installed

Contributions

- Unikernel images:
 - For reading sensor data
 - For calling external packages
- A way to create large-scale virtual clusters
- A hybrid Kubernetes cluster with cloud VMs, local VMs and IoT devices
- A Kubernetes application for managing virtual nodes (nodeWatcher)
- Virtual-kubelet fork that allows labeling at start
- Kubelet implementations:
 - A provider for managing Xen virtual machines
 - A provider for managing unix processes on Raspberry Pi

References

- A. Yenel. “Distributed Execution of Unikernel Applications on Container Orchestration Platform Kubernetes for IoT Scenarios”. Masterarbeit. Technische Universität München, 2020. ([mediaTUM](#))